


COMPUTER SYSTEM ORGANISATION

Naresh Jotwani

PowerPoint Slides



PROPRIETARY MATERIAL. © 2010 The McGraw-Hill Companies, Inc. All rights reserved. No part of this PowerPoint slide may be displayed, reproduced or distributed in any form or by any means, without the prior written permission of the publisher, or used beyond the limited distribution to teachers and educators permitted by McGraw-Hill for their individual course preparation. If you are a student using this PowerPoint slide, you are using it without permission.

Chapter 8

COMPUTER ARITHMETIC

Introduction

- Computer arithmetic is in principle very similar to usual decimal arithmetic.
- The main differences are:
 - (a) Computer circuits work on binary rather than decimal data
 - (b) Computer circuits are designed to operate at high speeds

Integer addition and subtraction:


In the process of digit-by-digit addition:

- *Carry out* is an ‘overflow’ from a digit position to the next higher digit position (on its left).
- *Carry out* from one digit position becomes *carry in* for the position to its left.
- If a digit position does not have *carry in*, we may set its *carry in* to 0.

These concepts apply to binary addition just as they do to the usual decimal addition.

- In the next figure, assume a_i , b_i and c_{i-1} as inputs, and s_i and c_i as outputs of the digital circuit shown.
- This circuit is known as a *full adder*.
- Inputs a_i , b_i and c_{i-1} will be either '0' or '1', and the circuit will add up their values.

One stage of integer addition in binary using a *full adder* -



- Between the input bits, there is no relative place value – i.e. ‘1’ has numerical value one, whether it occurs as the value of a_i , or b_i , or c_{i-1} .
- i.e. we need to simply count how many of these input bits are ‘1’.
- This count of ‘1’ bits on the three inputs is in the range from 0 to 3. We encode this count into two bits, which are the output bits c_i and s_i .
- There is relative place value between the output bits. Since carry out c_i is taken into the next position on the left, its place value is twice that of s_i .

- The next slide shows the complete truth table of the *full adder*.
- The table has $2^3 = 8$ rows.
- The first row corresponds to all inputs being '0', i.e. the total count of '1' bits is zero.
 - c_i and s_i remain '0'.
- The next three rows correspond to any one input to the circuit being '1'.
 - No *carry out* is generated in these three cases, and the outputs are $c_i = '0'$ and $s_i = '1'$.

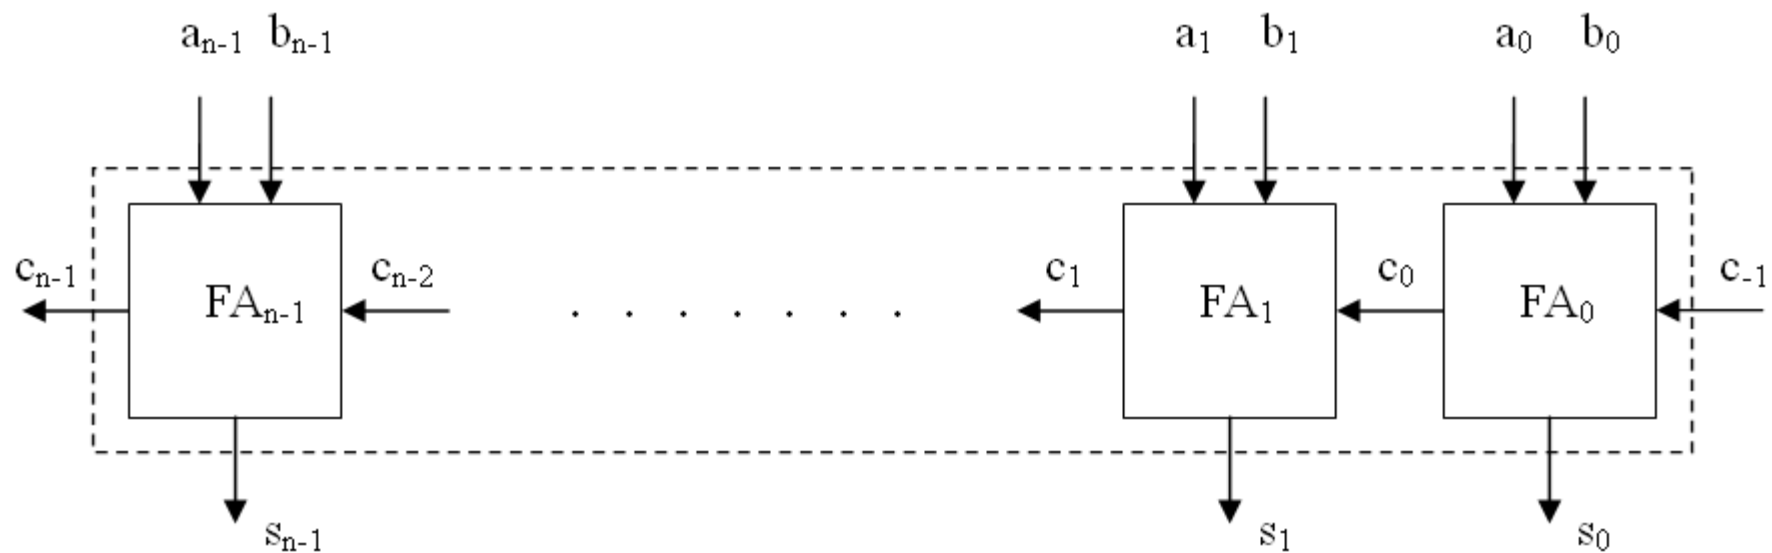
Truth table of full adder:

Inputs			Outputs	
A_i	b_i	c_{i-1}	c_i	s_i
0	0	0	0	0
0	0	1	0	1
0	1	0		
1	0	0		
1	1	0	1	0
1	0	1		
0	1	1		
1	1	1	1	1

- The next three rows correspond to any two inputs to the circuit being '1'.
 - A *carry out* is generated in these three cases, and the outputs are $c_i = '1'$ and $s_i = '0'$.
- The last row corresponds to all three inputs to the circuit being '1'.
 - A *carry out* is generated and the sum bit s_i becomes '1'. Outputs are $c_i = '1'$ and $s_i = '1'$.
- The circuit is a combinational circuit – the output at any time is a function only of the input at that time. Circuit has no memory of prior inputs or outputs.

- Logic gates can be used for a realization of the circuit, with maximum gate delay of 2.
- This *full adder* circuit can perform the work of binary addition at one bit position.
- In the next slide, n of these full adders are connected for full n -bit addition.
- The only ‘cause-and-effect’ relationship between the stages is:
 - Carry out c_i from bit position i becomes carry in c_{j-1} at bit position $j = i+1$ to its left

Circuit for n -bit addition, using n full adders



- We need to ‘string together’ n instances of the full adder, i.e. *carry out* bit of one stage becomes the *carry in* bit of the next stage to its left.
- The *carry in* bit of the rightmost stage – i.e. bit position 0 – has been numbered as c_{-1} .
- The *carry out* bit of the leftmost stage becomes the final *carry out* of the n -bit adder.
- At the end of addition, this bit is made available to programmer as the *carry* condition flag C.

- The dashed outer rectangle encloses the n full adders connected together. Known as *ripple carry adder*.
- This complete circuit has as its inputs the two n -bit integers A and B , and the *carry in* c_{-1} shown on the right.
- Outputs of the circuit are the n -bit sum S , and *carry out* c_{n-1} shown on the left.
- But, although this circuit is fully functional, it has a shortcoming.

- In any digital circuit, there is a *time delay* between the time inputs are applied and the time when outputs become stable.
- In the n -bit adder, the maximum number of logic gates is encountered in the paths between inputs the a_0 , b_0 and c_{-1} on the right, and the outputs s_{n-1} and c_{n-1} on the left.
- These circuit paths pass through each of the n full adders in the circuit, from right to left.
- Each full adder on the path contributes 2 gate delays, and therefore the total gate delay of the circuit is $2n$, unacceptable from a performance point of view.

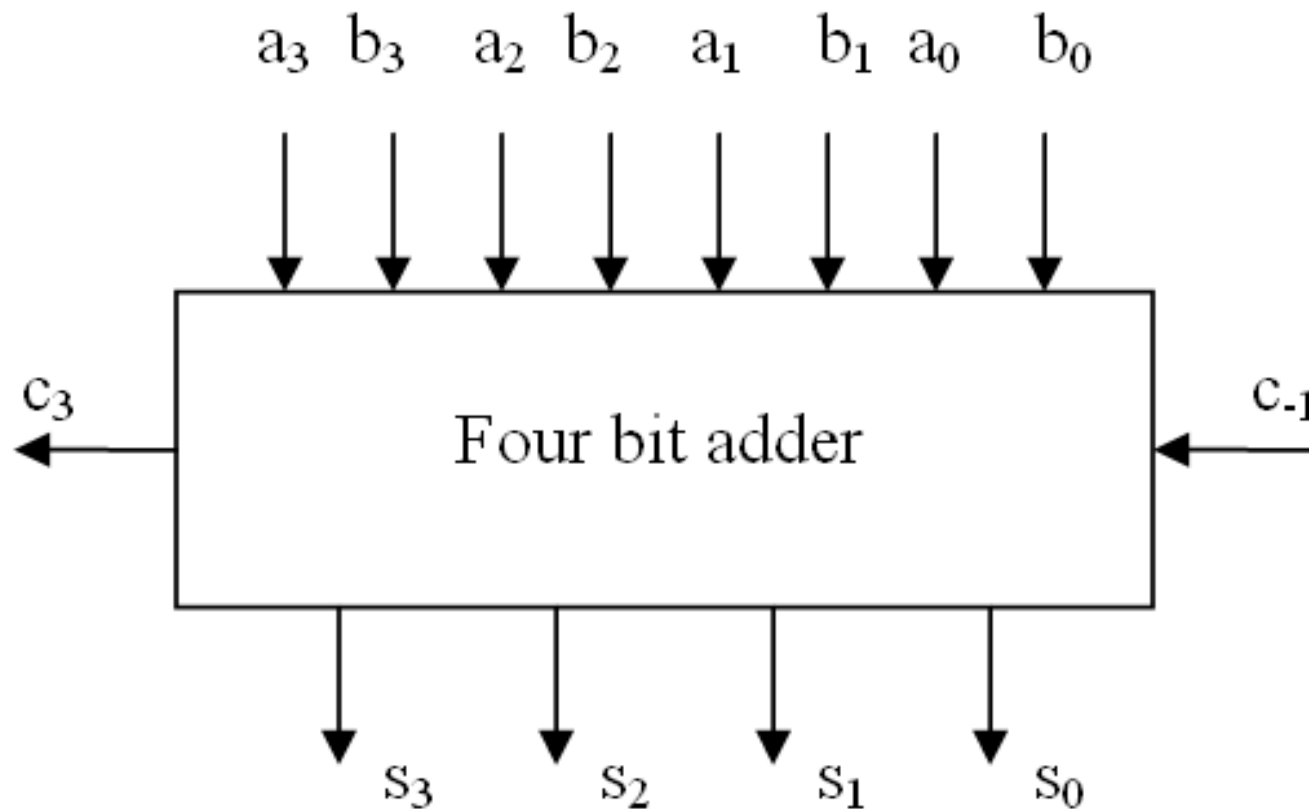
Carry Look-ahead:

For the n -bit adder:

- Suppose the building block is a 2-bit adder which also provides for *carry in* and *carry out*.
- Connect $n/2$ of such building blocks together to obtain the required n -bit adder. Carry ripples through these $n/2$ stages,.
- Each 2-bit adder can be designed with 2 gate delays
- The complete n -bit adder will now have $n/2 \times 2 = n$ gate delays
- Faster by a factor of 2, as compared to the previous *ripple carry* circuit with $2n$ gate delays

- We can also use 4-bit adders.
 - $n/4$ of the 4-bit adders can be connected for an n -bit adder. with a total gate delay of $n/4 \times 2 = n/2$.
- The next slide shows such a 4-bit adder.
- The two 4-bit inputs are $A = a_3a_2a_1a_0$ and $B = b_3b_2b_1b_0$; the 4-bit sum is $S = s_3s_2s_1s_0$
- *Carry in* is c_{-1} and *carry out* is c_3
- This 4-bit adder can also be thought of as a single digit hexadecimal adder.

A 4-bit adder as a combinational building block:



- *Carry look-ahead* is a systematic technique for designing multiple bit adders.
- Consider the 1-bit full adder discussed earlier.
- In the i^{th} stage of an n -bit adder, we define the following two functions of input bits a_i and b_i :

Generate function $g_i = a_i \cdot b_i$

Propagate function $p_i = a_i + b_i$

- Then *carry out* bit c_i is given by:

$$c_{i+1} = g_i + p_i \cdot c_{i-1}$$

- Generate bit g_i indicates that the i^{th} stage *generates carry out*, regardless of whether there is carry coming into this stage.
- Propagate bit p_i indicates that the i^{th} stage *propagates* incoming carry, regardless of whether there is carry generated in this stage.
- In stage i , *carry out* c_i is '1' if:
 - (a) stage i generates a carry, AND/ OR
 - (b) stage i receives incoming carry c_{i-1} , which it propagates to the stage on its left
- For the left-most stage of a 4-bit adder, we write:

$$c_3 = g_3 + p_3 \cdot c_2$$


- Substituting in this the expression for c_2 :

$$c_3 = g_3 + p_3 \cdot (g_2 + p_2 \cdot c_1)$$

- By substituting for c_1 , then for c_0 in the resulting expression, we obtain an expression for c_3 in terms of *carry in* c_{-1} and the input bits a_i and b_i , $i = 0$ to 3 .
- THUS \rightarrow Carry bits c_i , for $i = 0$ to 3 , depend on input bits only, and we do not require a carry to ripple forward stage by stage from right to left.
- Sum bit s_i of each stage is a function of a_i , b_i and c_{i-1} , as before in the 1-bit full adder of the i^{th} stage.
- The slow ‘ripple carry’ from right to left has been eliminated.

- In this way, a 16-bit adder can be built by connecting together four of the 4-bit adders, which are built using carry look-ahead.
- But, if this is done, a carry still has to ripple through from right to left, from one 4-bit adder to its left neighbour.
 - This is because, while each 4-bit adder internally makes use of carry look-ahead, there is no carry look-ahead between stages.
- Fortunately, as shown below, the principle of carry look-ahead also applies to a set of 4-bit adders connected together.

- For this, *generate* and *propagate* functions must be calculated for each 4-bit adder, and utilized to avoid the carry ripple.



- In the diagram, notation ‘g’ and ‘p’ indicates *generate* and *propagate* functions, respectively, from the single 4-bit adder shown.
- Carry look-ahead logic is shown above it, as a separate sub-circuit, to which four such 4-bit adders are connected.

Unsigned integer multiplication:

- *Repeated addition* proves to be a very inefficient method for larger integers.
- *multiply-shift-and-add* method is faster and more efficient. It involves multiplication by one multiplier digit at a time, as illustrated below:

$$\begin{array}{r}
 213 \\
 \times 379 \\
 \hline
 1917 \\
 14910 \\
 63900 \\
 \hline
 80727
 \end{array}
 \quad \left. \vphantom{\begin{array}{r} 1917 \\ 14910 \\ 63900 \end{array}} \right\} \begin{array}{l} \text{Partial} \\ \text{products} \end{array}$$

- In the second partial product, note that the number $7 \times 213 = 1491$ is shifted one position to the left.
- This is because this shift multiplies 7×213 by 10, to give us the required partial product 70×213 .
- The third partial product involves a left shift by two digits.
- Thus the multiplier 379 is broken up as $9 + 70 + 300$, and each partial product represents the product of one of these terms with the multiplicand 213.
- Multiplication of two n digit integers will in general produce a product which has $2n$ digits.

- *Multiply-shift-and-add* is an efficient method of multiplication; multiplication of any two n digit numbers requires the addition of n partial products.
- In binary, this is the method employed in computer systems for integer multiplication.
- *Multiply-shift-and-add* method applied to binary digits is shown below:
 - Let the four bits of the multiplier 1011 be denoted by $b_3b_2b_1b_0$ respectively.

$$\begin{array}{r}
 1 1 0 0 \\
 \times 1 0 1 1 \\
 \hline
 1 1 0 0 \\
 1 1 0 0 \mathbf{0} \\
 0 0 0 0 \mathbf{0} \mathbf{0} \\
 1 1 0 0 \mathbf{0} \mathbf{0} \mathbf{0} \\
 \hline
 1 0 0 0 1 0 0
 \end{array}
 \quad \left. \begin{array}{l} \\ \\ \\ \end{array} \right\} \begin{array}{l} \text{Partial} \\ \text{products} \end{array}$$

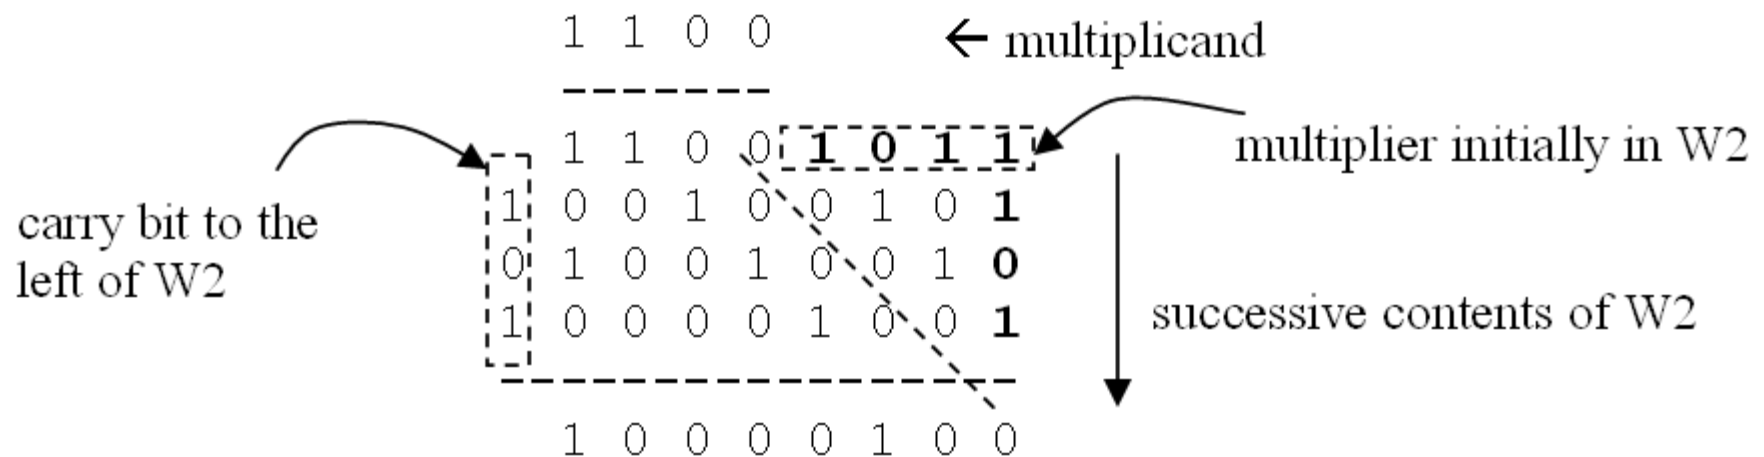
- Observe that:
 - (a) When a multiplier bit b_i is '0', the corresponding partial product is zero, and
 - (b) When multiplier bit b_i is '1', the corresponding partial product is simply a shifted version of the multiplicand 1100.
- Bold-face zeros in partial products show number of positions the multiplicand has been shifted to the left.

- For implementation on a computer:
 - (i) We reserve a result register, say W2, of size $2n$ bits.
 - (ii) To keep the multiplicand unchanged through the multiplication process, we shift W2 to the right at each step, instead of shifting the multiplicand to the left.

1	1	0	0					← multiplicand	
×	1	0	1	1					← multiplier
<hr style="border-top: 1px dashed black;"/>									
1	1	0	0	0					<div style="display: flex; align-items: center;"> <div style="font-size: 3em; margin-right: 10px;">}</div> <div>Partial products</div> </div>
1	1	0	0	0	0				
0	0	0	0	0	0	0			
1	1	0	0	0	0	0	0		
<hr style="border-top: 1px dashed black;"/>									
1	0	0	0	0	0	1	0	0	← Register W2, shifted right at each step

- The dashed inclined line above corresponds to the rightmost bit position of $W2$.
- Note that we do not need separate registers for the n partial products.
- Of course, we do need two separate n -bit registers for the multiplicand and multiplier.
- And we need the $2n$ -bit register $W2$ into which the partial products get added, and which stores the result.
- Note also that:
 $W2$ is initially zero, and only n of its $2n$ bits are utilized in the first *shift-and-add* step carried out.
Once a bit b_i of the multiplier has been examined, it is not needed again during the multiplication process.

- Thus the right half of W2 can be used to initially hold the multiplier. As a bit of the multiplier gets shifted out on the right, one more bit of W2 is available to store the accumulating product.
- Carry is a part of the accumulating product, and must be shifted in from the left during the next right shift.
- This arrangement is shown below:



- The four bits 1101 in bold-face shown on the right are successive bits of the multiplier.
- One by one they get shifted out. At any step, the multiplicand is added into W2 only if the particular multiplier bit is '1'.
- The inclined dashed line shows how the rightmost '0' bit of the multiplicand gets shifted right one bit position at a time, to finally become the rightmost bit of W2.
- The next slide shows this technique in algorithmic notation. The inner loop of the algorithm is repeated n times, for n -bit multiplication.

Algorithm for unsigned integer multiplication

```
initialize W2 (L) to zero, and
           W2 (R) to multiplier
multiplicand is in register M
repeat n times
{
    if rightmost bit of W2 is '1' then
        add M to W2 (L) , saving carry
    shift carry & W2 one bit to right
}
final product is in W2
```


- In a typical processor, the $2n$ -bit register W2 will be a pair of n bit general purpose registers, e.g. R3 and R4, and M a third such register.
- (a) ADD and SHIFT can be used to provide the multiply function in software.

- (b) For a hardware integer multiplier, the algorithm above must be translated into an appropriate *sequential* digital circuit, and provided as a functional unit.
- Integer multiply will then be a machine instruction, similar to ADD and SUB.
- (c) If the processor does not provide an integer multiply instruction, then it is possible to implement the multiply algorithm in *microcode*. Multiply instruction is thereby added to the instruction set.
- Microcoded implementation provides execution time faster than software, but slower than that of a hardware multiplier.

Booth's algorithm and integer division

- The previous algorithm cannot be used directly for multiplying signed integers.
- *Booth's algorithm* provides a faster alternative which also takes care of signed multiplication in a natural way.
 - Consider 8-bit multiplier 00111111, which equals decimal 63. This can be written as $64 - 1$, which in binary is: 01000000 - 00000001.
 - The first term represents a left shift of the multiplicand by 6 bit positions, whereby it is multiplied by 64.
 - The second term represents a subtraction of the multiplicand.

- Now consider the multiplier 00111100, which has decimal value 60. This can be written in binary as 01000000 – 00000100.
- The second term now represents left shift by two positions and a subtraction.
- In general, suppose the multiplier has a k consecutive '1' bits, within its total size of n bits. Then the relevant part of the multiplier has the form:




- Let the rightmost of these k '1' bits be b_j . Then the leftmost of these k bits is b_{j+k-1} . The numerical value of these k bits is given by:

$$\sum_{i=j}^{j+k-1} 2^i$$

- This is a summation of k different powers of 2, which would need k steps of *shift-and-add*.

- But the summation equals $2^{j+k} - 2^j$, which involves only one addition and one subtraction.
- If $k > 2$, then instead of k additions, we get by with one addition and one subtraction, for a total of only 2 operations.
- This is the essence of the Booth's algorithm for multiplication.

- The algorithm has the additional benefit that it naturally takes care of signed operands in 2's complement form.
- To see that, consider the following multiplier:



- The k '1' bits are now the leftmost k bits in the n -bit multiplier, i.e. $j+k-1 = n-1$. Now we get the equivalent value of the multiplier as: $2^{j+k} - 2^j = 2^n - 2^j$
- Value 2^n corresponds to bit b_n – which is not within the n bits $b_{n-1}..b_0$ of the multiplier. Without this bit, the value of summation reduces to -2^j , which is the numerical value of the leftmost k bits of the multiplier in 2's complement.
- Thus Booth's algorithm has the following advantages:
 - (a) sequences of '1' bits in multiplier are handled more efficiently
 - (b) signed operands are handled in a natural manner, in 2's complement representation.

Integer division:

- (i) integer division can be viewed as repeated subtraction;
 - (ii) in binary, the ‘long-hand’ method of division amounts to a sequence of *shift-and-subtract* operations.
- *Shift-and-subtract* steps of division can generate the integer division algorithm.
- It is possible to implement the division algorithm in software, hardware, or microcode.

Floating point operations:

- To add two floating point numbers, it is necessary to make their exponents equal, before the mantissas can be added.
- A floating point number in binary remains unchanged in value if its mantissa is shifted right by k bit positions and the exponent is increased by k .
- By applying this transformation to the operand with the smaller exponent, we can make the two exponents equal. Then the two mantissas can be added.
- The sum will in general be un-normalized, and must be normalized as the final step of the operation.

- Addition & subtraction of floating point numbers go through the same series of steps. Sign bits of the two operands must be taken into account during the process.
- For multiplication and division, the two respective exponents need not be equal.
- In multiplication, exponents of the two operands are added; in division, exponent of the divisor is subtracted from that of the dividend.
- **Example:** See floating point operation $32 - 15.5$ on next slide.

$$\begin{array}{rcl}
 + \text{ 32:} & 0 & 10000100 \quad 000000000000000000000000 \\
 - \text{ 15.5:} & 1 & 10000010 \quad 111100000000000000000000
 \end{array}$$

- We un-normalize the second number, i.e. shift its mantissa to the right by 2 bits, and increase its exponent by 2, to give:

$$1 \quad 10000100 \quad \underline{011111000000000000000000000000}$$

- Note there is no longer a hidden '1' bit, since the former hidden bit is now included in the $23+2 = 25$ bits of this mantissa. Two numbers have opposite signs, so we perform mantissa subtraction:

$$\begin{array}{r}
 \underline{1}. \quad 0000000000 \quad . \quad . \quad . \\
 - \quad 0. \quad \underline{011111000} \quad . \quad . \quad . \\
 \hline
 = \quad 0. \quad 100001000 \quad . \quad . \quad . \quad \leftarrow \text{result mantissa, to be re-normalized}
 \end{array}$$

Now we normalize the result, giving:

$$0 \quad 10000011 \quad 000010000000000000000000 = 16.5 \quad (\text{correct result})$$

- Three options for implementing floating point operations on a system: *software*, *dedicated hardware*, or *microcode*.
- On many processors, floating point operations are provided using a dedicated hardware unit known as *floating point unit (FPU)*.
- The two functional units ALU and FPU are usually capable of operating *in parallel*.
- Such parallelism of operations is a major potential source of performance enhancement.

Binary Coded Decimal (BCD) Arithmetic:

- Packed* BCD format uses a nibble – i.e. 4 bits – for each decimal digit. For example, the eight digit number 31052008 would be packed into four bytes as:

0011 0001	0000 0101	0010 0000	0000 1000
-----------	-----------	-----------	-----------

- On some processors, special machine instructions operate on numbers in packed BCD format. ‘Auxiliary carry’ must be provided from right nibble to left.
- Programming languages in earlier years provided a packed BCD type of number representation.

Conditions, exceptions, and faults:

- Condition flags are needed to provide required conditional transfers of control within a program.
- Instruction set of the processor provides for conditional transfers based on these flags. [Recall *NICE* instruction set seen in Chapters 4 & 5.]
- These condition flags do not represent ‘*errors*’, since they occur routinely in any running program.
- Programmer responds to these flags according to program requirements. Branch decisions made on that basis make up the *flow of control* within a program.

- But other conditions can arise in a running program which need to be handled in a different manner.
- *Floating point overflow, underflow, divide by zero* etc. are examples of such conditions.
- Such conditions do not cause any condition flags to be set in PSW, but they are treated as ‘run-time’ errors, and they also known as *exceptions*.
- Yet another type of condition which can arise within a processor is termed as a *fault*.
- Examples: *memory access violation, page fault* [to be studied in Chapter 9].

- In such cases, it is not possible for the program itself to take corrective action.
 - Rather, exceptions and faults are handled by *interrupts*. The source of an interrupt may be:
 - (a) an external input or output device, such as keyboard or disk [to be studied in Chapter 10],
 - (b) An exception or fault occurring within the processor, such as floating point overflow or memory violation,
 - (c) a *software interrupt*, which is a machine instruction executed by the running program.
- [Note: At present, *NICE* instruction set does not provide for software interrupt].

Summary

- Full adder, n -bit ripple carry adder
- Use of carry look-ahead technique
- Shift-and-add technique for multiplication
- Booth's algorithm
- Floating point operations
- Conditions, exceptions, and faults encountered in a running program